

# Welcome

This is the source code for CodFS described in our paper presented in USENIX FAST 2014. The system is tested on Ubuntu 12.04.2 64-bit with GCC 4.6.3.

Release 1.0.1 (20-6-2014) Code cleanup and documentation update

## Installation

These are the required libraries that users need to download separately. Brackets denote the package names in Debian and Ubuntu platforms. Users can use apt-get to install the required libraries.

- build-essential (build-essential)
- scons (scons)
- boost libraries version 1.5x (libboost-dev, libboost-program-options-dev, libboost-thread-dev, libboost-filesystem-dev)
- FUSE (libfuse-dev)
- openssl (libssl-dev)
- pkg-config (pkg-config)

For Debian and Ubuntu users:

```
sudo apt-get install build-essential scons libboost-dev libboost-program-options-dev libboost-t  
hread-dev libboost-filesystem-dev libfuse-dev libssl-dev pkg-config
```

The following libraries have to be compiled and installed manually:

### Google Protocol Buffers

```
$ cd lib/protobuf-2.4.1  
$ ./configure; make clean; make  
$ sudo make install  
$ sudo ln -s /usr/local/lib/libprotobuf.so.7 /usr/lib/libprotobuf.so.7  
$ sudo ln -s /usr/local/lib/libprotoc.so.7 /usr/lib/libprotoc.so.7
```

### MongoDB C++ Driver (only required for MDS)

```
$ cd lib/mongo  
$ sudo scons install
```

## Preparation

### Compile

The program can be compiled using Linux make. Running make on the CodFS root will automatically compile all CodFS components.

```
$ make
```

For optimal performance, turn off debug mode by editing the following flags in Makefile:

```
OPTIMIZE := -O3  
EXTRA_CFLAGS := -std=c++0x -DDEBUG=0 -D_FILE_OFFSET_BITS=64
```

### Setting up MongoDB

We suggest hosting the MongoDB on the MDS for optimal performance.

- Set up MongoDB environment

On the MongoDB host, create the database parent directory

```
$ sudo mkdir -p /data/db
```

Change the MongoDB address "[mongodbip]:[mongodbport]" in "mongo-codfs.js" if necessary

```
db = connect("localhost:27017/codfs");
```

Change the default MongoDB password in "mongo-codfs.js" for security reasons

```
"pwd": "cf8b7fa0b1bcd277da8217f04f498bd0"
```

#### Note

- "pwd" should be an MD5 hash, which can be generated by running `echo 'NEWPASSWORD' | md5sum` in Linux
- Change to the directory containing the MongoDB binary

```
$ cd mongo/bin
```
- Start MongoDB daemon

```
$ sudo ./mongod --fork --dbpath /data/db --logpath /var/log/mongodb.log --logappend
```
- Import MongoDB Setting

```
$ ./mongo ../../mongo-codfs.js
```

## Adjusting XML Settings

- Copy one of the two set of example XML files from the directory to the CodFS root (same level as the executables) For standalone setup, cp `example_xml/standalone/*` . For distributed setup, cp `example_xml/distributed/*` .
  - `example_xml/standalone/`
    - Example XML files for a standalone setup which runs all components on localhost (127.0.0.1)
  - `example_xml/distributed/`
    - Example XML files for a distributed setup
      - MONITOR on 192.168.0.11:53000
      - MDS on 192.168.0.12:50000
      - MongoDB daemon on 192.168.0.12:27017
- Adjust the following essential parameters in the XML files according to the cluster setup:
  - `mdsconfig.xml`
    - MDS listen port
    - MongoDB IP address, listen port, password
  - `osdconfig.xml`
    - Update scheme (FO=0, FL=1, PL=2, PLR=3) and reserved space size
  - `monitorconfig.xml`
    - MONITOR listen port
  - `clientconfig.xml`
    - Coding scheme (e.g., RS, RDP) and coding settings (e.g. n, k)
  - `common.xml`
    - MDS IP address and listen port
    - MONITOR IP address and listen port
- Refer to comments in the XML files for other individual settings (e.g., OSD capacity)

## Running CodFS

### Server-side

Copy the corresponding executable (MONITOR, MDS or OSD) and all the XML files to each server

Start the components in the following order

- MONITOR

```
$ ./MONITOR
```
- MDS

```
$ ./MDS
```
- Start OSD one-by-one

```
$ ./OSD [component_id] [network_interface]
```

#### Note

- Assign a `component_id` that is unique from MDS, MONITOR, other OSDs and clients
- Use "ifconfig" to check the name of the network interface (e.g., eth0). If multiple interfaces exist, select the one that can be reached by other components

## Client-side

On the client machine, create two directories

```
$ mkdir fusedir mountdir
```

Mount the FUSE volume with the following options

```
$ ./CLIENT_FUSE -o big_writes,large_read,noatime -f mountdir [component_id]
```

You can now access CodFS through `mountdir`. The following command copies a file into the cluster and outputs the time spent:

```
$ time cp testfile mountdir/
```

#### Note

- Assign a `component_id` that is unique from MDS, MONITOR, OSDs and other clients

## Recovery

### Automatic Recovery

CodFS does not automatically recover by default. To enable this features, uncomment the line `#define TRIGGER_RECOVERY` in `src/common/define.hh`

#### Note

- Adjust the sensitivity of recovery trigger by changing `SleepPeriod`, `DeadPeriod` and `UpdatePeriod` in `monitorconfig.xml`

### Manual Recovery

Manual recovery is useful in benchmarks.

To simulate failure, kill the OSD process on one of the OSD hosts

```
$ pkill -9 OSD
```

To force failure detection and recovery, send a signal to the MONITOR process on the MONITOR host

```
$ pkill -USR1 MONITOR
```

## Benchmarks

### Test seq. read/write throughput using dd

After mounting the FUSE volume, test seq. write throughput by:

```
$ dd if=/dev/zero of=mountdir/ddtest count=256 bs=16M conv=fdatasync
```

Drop page cache:

```
$ sync && sudo echo 3 > /proc/sys/vm/drop_caches
```

Test seq. read throughput by:

```
$ dd if=mountdir/ddtest of=/dev/null count=256 bs=16M
```

## Test random write throughput using IOzone

Download latest IOzone tarball from <http://www.iozone.org/>.

Compile IOzone from source

```
$ tar xf iozone3_424.tar
$ cd iozone3_424/src/current
$ make linux-AMD64
```

Copy the IOzone executable to CodFS root, and execute the following:

```
$ ./iozone -Ra -+n -i0 -i2 -s 4g -r 128k -e -c -w -f mountdir/iozonetest
```

This will create a 4GB file and perform random read/write with 128KB record size.

## Contact

Patrick P. C. Lee (<http://www.cse.cuhk.edu.hk/~pclee>)

---

Formatted using GitHub Flavored Markdown.